

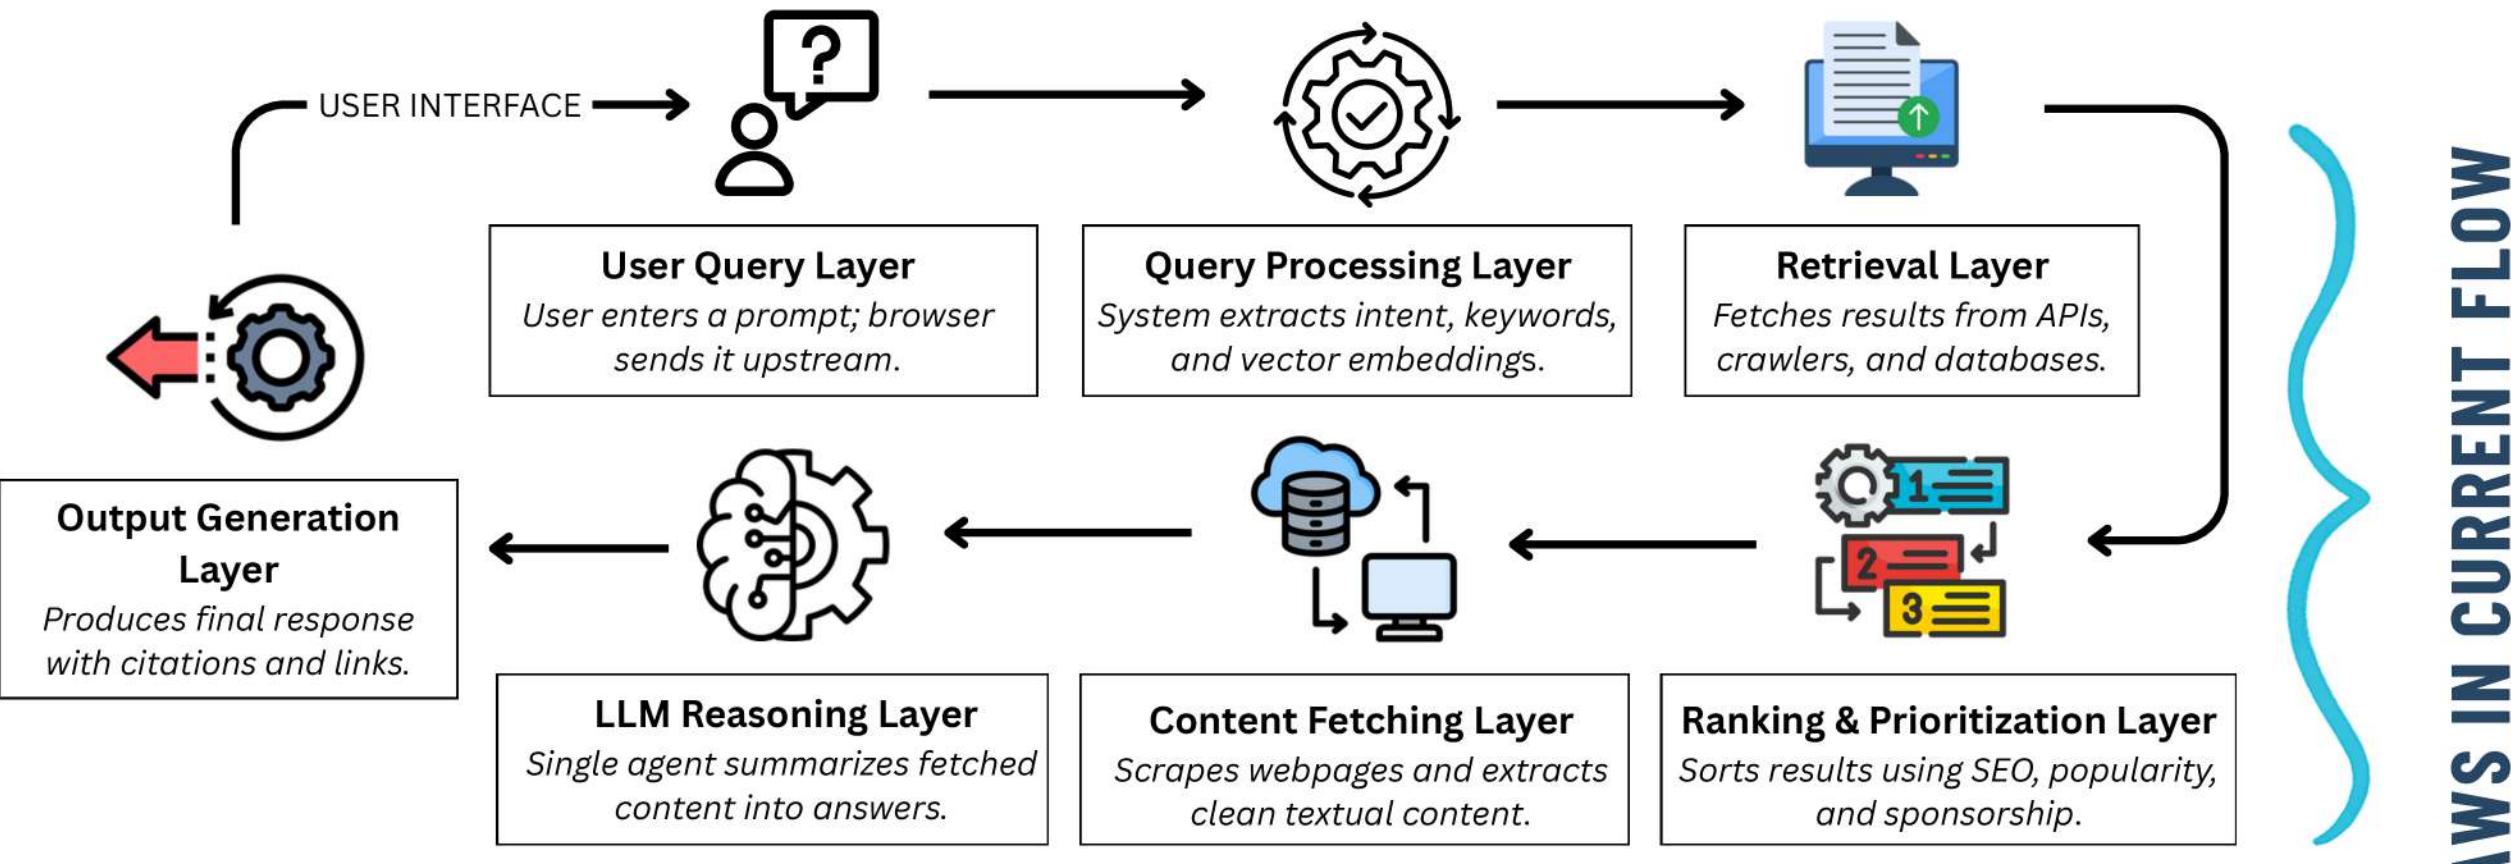
# NEXUS

## Multi-Agent Browser

**ADI AMRUTA**  
**IIT KANPUR**  
**ROLL NO.-220058**  
**CONTACT NO.-9861106538**



# TODAY'S AI BROWSING PIPELINE



## PROBLEMS EXISTING TODAY

### SEVERE PRODUCTIVITY WASTE

- Knowledge workers lose **2.5 hours** daily to unproductive digital tasks
- 600+** hours annually per worker wasted on information retrieval
- \$54,000+** lost productivity per worker per year (at \$90/hour)

#### Global Scale:

- \$1.8–1.9 trillion** in annual global productivity loss
- 350M** knowledge workers affected worldwide

### BROKEN INFORMATION ECOSYSTEM

- 60%+** of top results are ads or sponsored content (not merit-based)
- SEO gaming** pushes authoritative sources below optimized blogs
- AI search **amplifies bias** by surfacing "accessible but low-rigor content"

#### The Authority Problem:

- Paywalled academic sources get **preferential AI training data access**
- Corporate-sponsored** research outranks independent analysis

### CONTEXT & WORKFLOW CHAOS

- Average professional has **14-36 tabs** open simultaneously
- 45 context** switches per hour between apps and tabs
- 23 minutes** required to regain deep focus after each interruption

#### Repetitive Work Epidemic:

- 64%** of workers repeat the same research tasks monthly.
- Zero memory between sessions—every search starts from scratch
- 5.3 hours** weekly **lost** to manual work

Sources: McKinsey Global Institute 2023, IDC Future of Work 2024, Forrester Research

Nature 2023, Pew Research Center, Stanford Internet Observatory

RescueTime 2024, Atlassian State of Teams, Microsoft Work Trend Index, Forrester

### Browsers Not Intelligent – (Step 1-2)

- Browsers don't understand what the user is actually trying to accomplish – no grasp of **goals, depth, or relevance**.
- They simply take the query and display webpages, leaving all **planning and reasoning** entirely **on the user**.

### Search Prioritizes Money, Not Relevance – (Step 4)

- Ranking is heavily influenced by **ads, SEO strength, and domain authority** rather than true quality or rigor.
- As a **result**, **sponsored articles** and optimized content outrank deeper, more relevant insights.

### AI Assistants Are Limited – (Step 6)

- LLMs work in a **single threaded**, one-shot manner with no multi-agent collaboration or **long-term context**.
- They cannot break **complex research** into steps or remember previous work, **limiting their usefulness** for real workflows.

### No Unified Research & Workflow Automation (System-Wide)

- Today's research spans isolated apps with no shared state, **unified memory**, or **orchestration layer**, forcing every component of the workflow to run independently.
- browser lacks persistent **task-state retention**, and **workflow scheduling**, users must manually reconstruct their research state.

## ASSUMPTIONS

### 1. Market Assumptions

Enterprises pay \$30–\$200/user/month for productivity gains

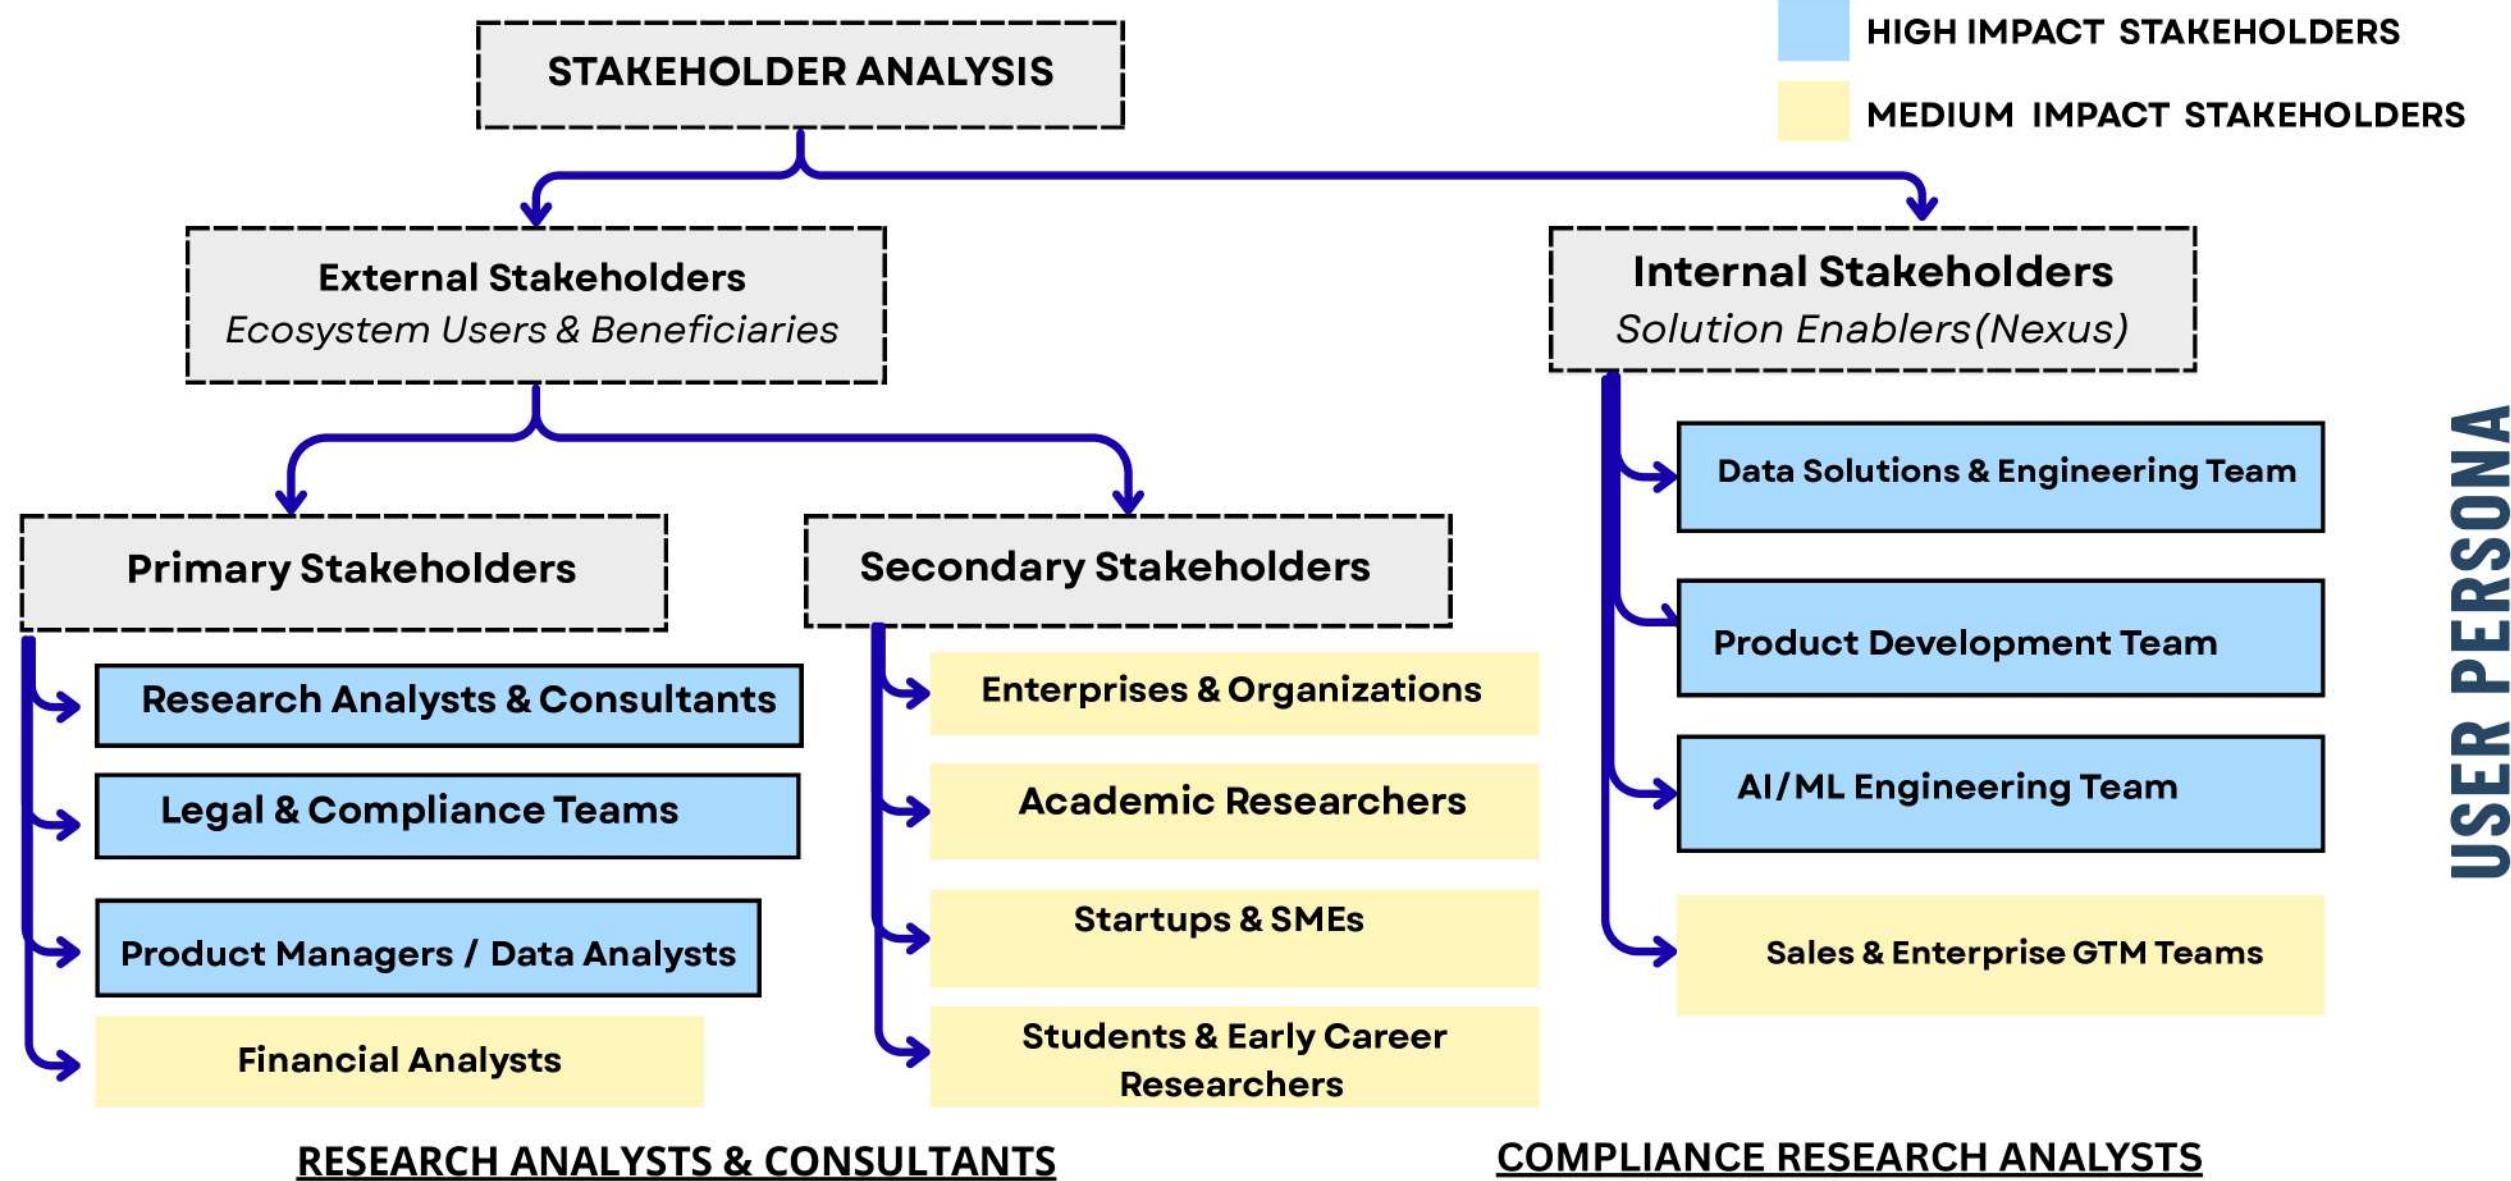
### 2. Technical Assumptions

Multi-agent outperforms single-agent systems on complex tasks  
Workers accept 30–90 secs latency in exchange for hours saved

### 3. Behavioral Assumptions

Users trust automation when sources are transparent

Persistent research memory creates strong switching costs

**PAIN POINTS**

- Spend 60-70% of time gathering and validating information instead of analyzing
- Constant tab overload across tools (Google, Insights, email)
- Repetitive research cycles with no automation or memory

**USER NEEDS**

- A multi-agent system to run research tasks in parallel
- A browser that remembers context, prior searches, sources project threads
- Automated synthesis, cross-validation, insight extraction

**GOALS**

- Reduce research time from 6 hours → 20-30 minutes
- Deliver data-backed insights to senior consultants faster

**PAIN POINTS**

- Must track regulatory updates manually across portals
- Risk of errors due to manual verification.
- Relevant legal content is buried under SEO-heavy sites.

**USER NEEDS**

- Automated cross-validation and monitoring of official regulatory sources.
- A secure browser that prioritizes official documents and maintains audit trails.
- Tools for quick comparison, clause extraction, and change detection.

**GOALS**

- Achieve zero compliance breaches with accurate, verifiable updates.

**USER PERSONA****NAME:** Aarav Khanna, AGE: 27**ROLE:** Senior Research Analyst, Consulting Firm**LOCATION:** Bangalore, India**WORK CONTEXT:** High-intensity research cycles, tight deadlines, client-facing deliverables**Goals**

- Produce faster, deeper insights for partners
- Reduce manual work (data gathering, summarizing, valid)
- Deliver high-quality slides backed by accurate data
- Build a repeatable research system for recurring industries

**Pain Points**

- Loses hours daily in tab overload and context switching
- Struggles to keep track of sources, citations, evidence
- AI tools provide answers but don't execute multi-step workflows
- Stuck validating which pages are sponsored vs. truly relevant

**What Aarav Needs from NEXUS ?**

- A multi-agent system that performs parallel research
- Automated synthesis of market, competitor, and funding data
- Auto-verification across multiple independent sources
- A browser that remembers past research and reuses workflows
- Non-sponsored, relevance-first ranking

**NEXUS INTERNAL TEAM****PAIN POINTS**

- Must orchestrate multi-agent workflows reliably across diverse, unstructured web data.
- Need to manage accuracy, bias detection, and verification at scale while keeping latency low.

**USER NEEDS**

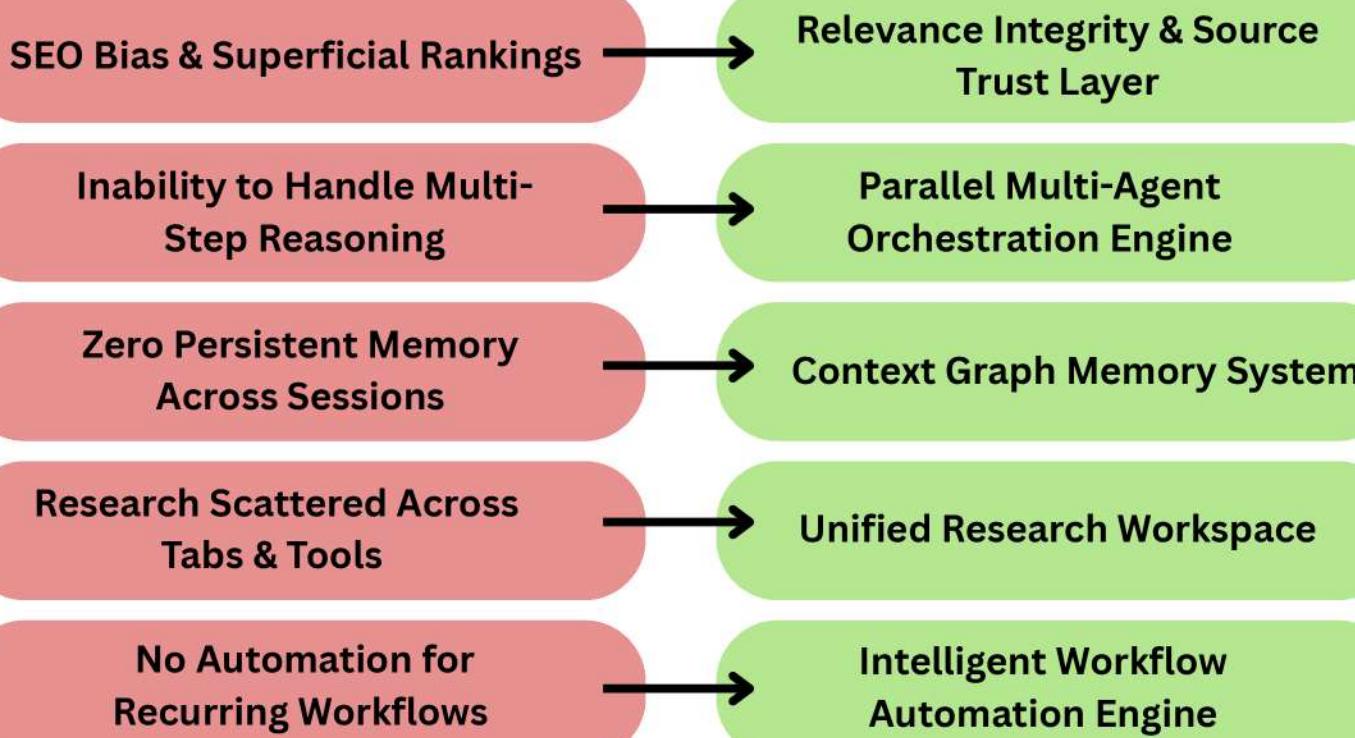
- A scalable agent orchestration framework with persistent memory and task-state management.
- Robust monitoring tools for agent performance, workflow failures, and system health.

**GOALS**

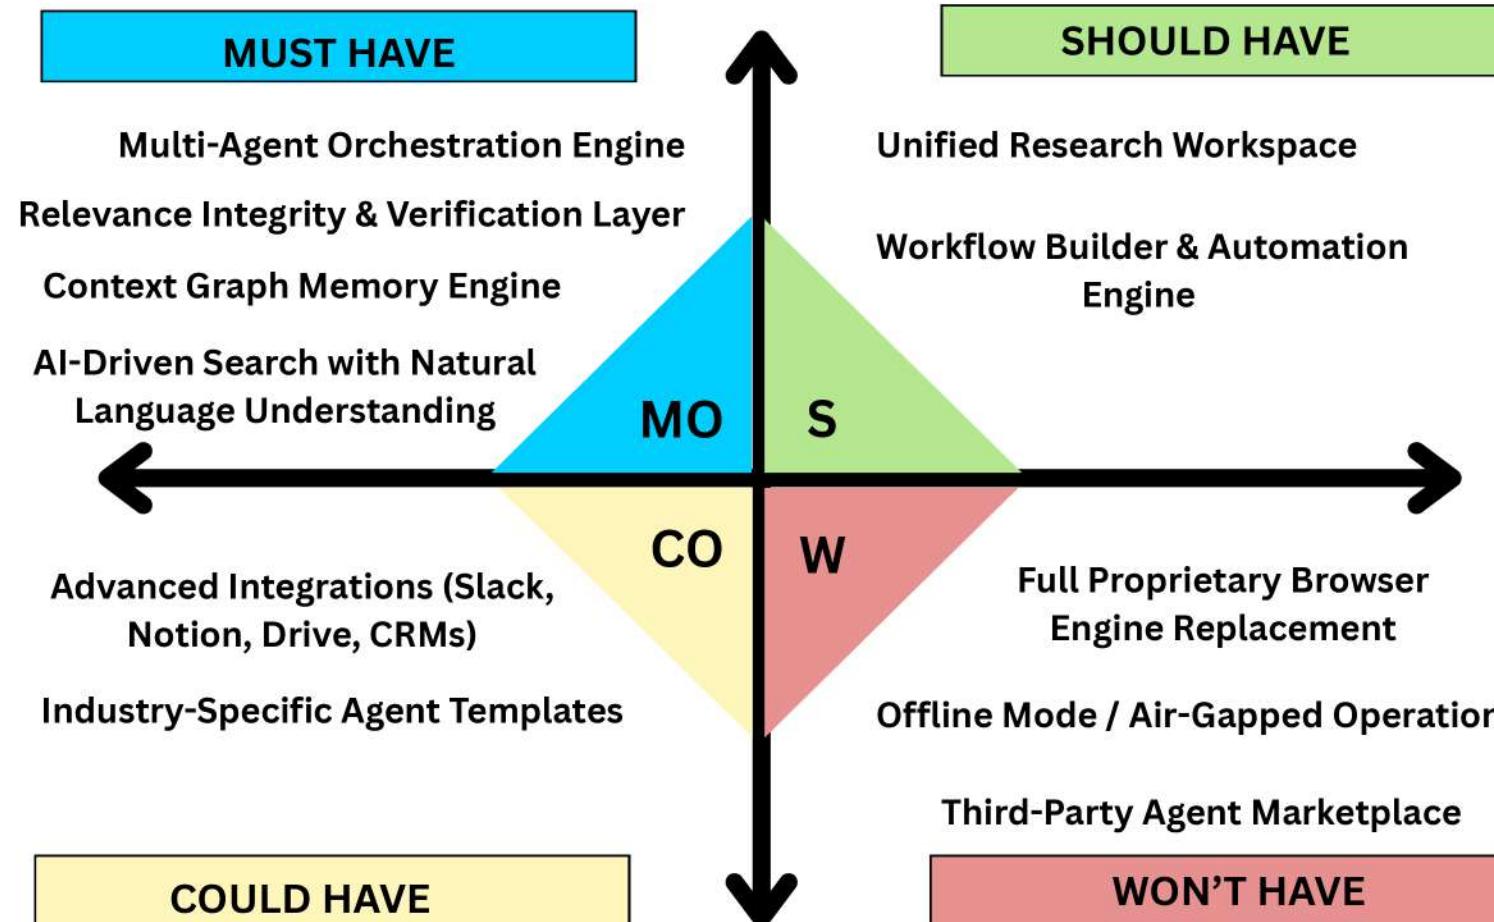
- Deliver a stable, accurate browser with 10x user productivity.



# ⚠ PROBLEM → SOLUTION Mapping

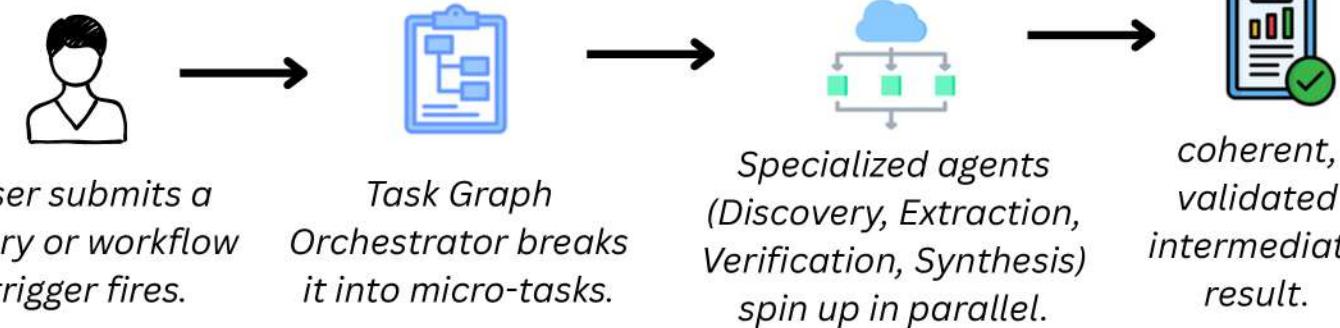


## NEXUS MVP – Feature Prioritization



## SOLUTION ARCHITECTURE (NEXUS 4-Layer Architecture)

### LAYER 1 : AGENT INTELLIGENCE LAYER



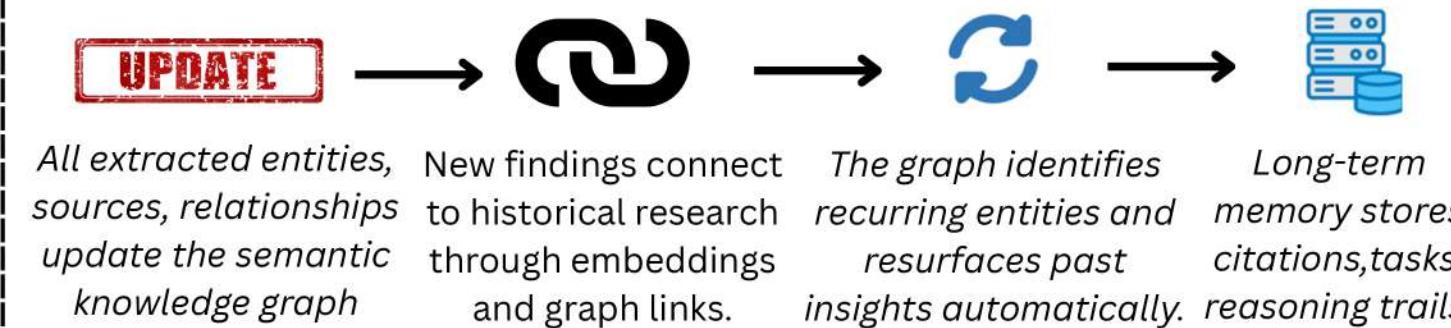
### What NEXUS does differently?

- Prioritizes **official sources**, primary research data
  - Runs a **source integrity score** (SIS) to filter out biased pages
- Result: NEXUS delivers truth-first research, not SEO-first data

### LAYER 2 – RETRIEVAL & VERIFICATION LAYER



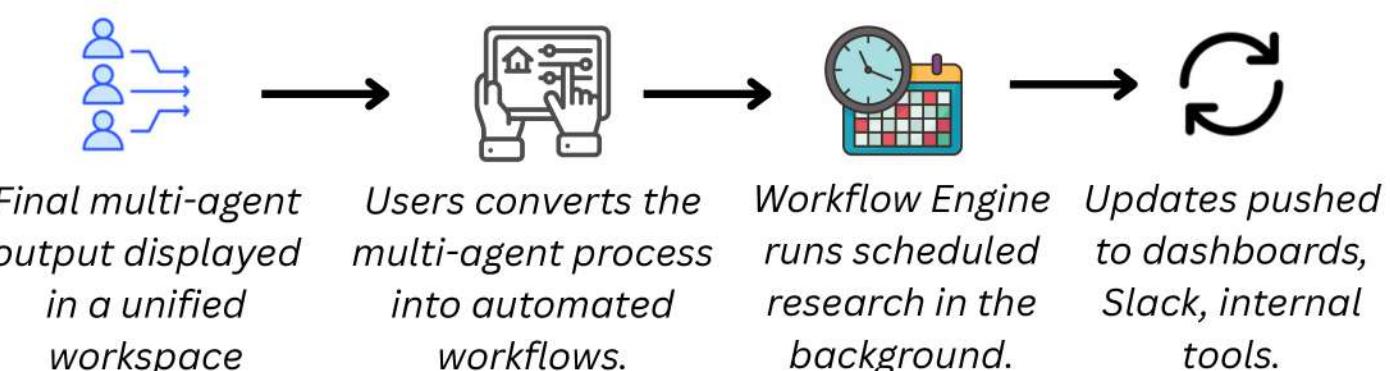
### LAYER 3 – MEMORY & RESEARCH GRAPH LAYER



### What NEXUS does differently?

- Maintains **persistent, session-independent** memory
  - Resurfaces **past insights** using a semantic knowledge graph.
- Result: Analysts don't start from zero every time

### LAYER 4 – WORKFLOW & WORKSPACE LAYER



### What NEXUS does differently?

- Creates a single workspace for all research tasks
  - Agents are set to monitor, alert, and update automatically
- Result: No repeated work, no manual tracking.



# UNIFIED ADOPTION STRATEGY

*"NEXUS is not positioned as a browser, it is a operating system built for high-stakes knowledge work."*

## Stage 1 – Vertical Beachhead (Analysts, Consultants, Finance)



NEXUS enters through the users who feel the highest research pain and experience the clearest time-to-value.

### Why this works ?

- These users spend **60-70% of time** synthesizing information
- Strong **peer networks** → (+ve within group effect)
- Extremely high willingness to adopt productivity tools

## Stage 2 – Team & Department Expansion (Horizontal Spread)



Once a few individuals adopt NEXUS, team expansion becomes natural due to Coordination Lift.

### Why this works ?

- Shared workflows → reusable processes
- Shared memory graphs → shared context
- Shared entities → team-level knowledge alignment

## Stage 3 – Enterprise Rollout (Scale Across Org)



As organizations accumulate workflows, insights, and memory graph intelligence, switching away becomes extremely costly.

### Why this works ?

- More research → richer memory → smarter agent recommendations → more reliance.
- This creates structural, compounding switching costs.



## HOW NEXUS WINS EARLY?

- Industry-specific agent templates (“Funding Deep Dive”, “Market Landscape Scan”)
- Fast onboarding (first insight < 5 minutes)
- Case-study driven GTM (before/after comparisons)
- Outcome: Rapid traction, clear PMF signals, organic analyst-to-analyst adoption.

## WHAT MAKES ADOPTION INEVITABLE?

- Workflow Network Effects
- Compounding Context Graph Memory
- Daily Automation Hooks

- Daily Automation Hooks
- High Perceived Value vs. Low Switching Cost at Entry

## GROWTH TECHNIQUES (AARRR FRAMEWORK)

### Acquisition

- Browser extension + one-click onboarding
- LinkedIn content targeting analysts & PMs
- Organic demos from shared PDFs

### METRICS

- Visit → Install Con. Rate (15-25%)
- Install → Sign-Up Rate (35-50%)
- % Organic Referral Traffic (20-30%)

### Activation

- First insight under 2 minutes onboarding
- Prebuilt agent templates
- Auto-suggested workflows

### Retention

- Daily automated monitoring workflows
- Long-term memory graph
- Team collaboration features

### Referral

- Created with NEXUS in exported reports
- Internal workflow sharing

### Revenue

- Analyst community templates

- Activation Rate (Sign-Up → First Insight: 40-55%)
- Agent Template Usage in First Session ( $\geq 1$ )

- Day-7 Retention (35-45%)
- Day-30 Retention (20-30%)
- Weekly Automated Workflow Runs/User (5-12)

- Free → Pro Conversion (5-12%)
- Team Seat Expansion Rate (25-35%)
- Enterprise Seat Multiplier (2.5×)

- Internal Workflow Sharing (30-50% of teams)
- Workflow Marketplace Downloads/User (1-3 per week)



# GO-TO-MARKET STRATEGY

## Hyper-Targeted Vertical Wedge (High Pain → Fast Adoption)

Deploy Verticalized Workflow Templates

- Funding Deep Dive (company → investor → sector mapping)
- Competitive Tracker (multi-source ingestion + synthesis)

### Activation Hook: 3 Free Workflow Credits/Day

- Forces first experience of speed & accuracy
- Users quickly realize manual research is no longer acceptable



## Value in <120 Seconds (The Hook Moment)

Multi-Agent Parallel Execution

- Retrieval agent → gathers & ranks sources
- Verification agent → cross-checks 3+ sources
- Synthesis agent → builds insight graph → summary

### Insight Delivered in Under 2 Minutes

- User immediately feels the shift from “searching” to “knowing”
- This becomes their irreversible adoption moment



## Team-Level & Department-Level Growth (Coordination Lift)

### A. What It Does

- Enables shared workflows across teams
- Creates a unified context graph for collective knowledge
- Synchronizes dashboards and recurring research tasks

### B. How It Delivers Impact

- Team members reuse and refine each other’s workflows
- Shared memory reduces duplication of research
- 2-3 early adopters naturally pull in the entire department



## HIGH-LEVERAGE CONTENT ENGINE

- What It Does**
  - Publishes benchmark reports and templates
  - Operates a workflow/template marketplace for users
- How It Delivers Impact**
  - Establishes NEXUS as the authority in AI research workflows
  - Generates steady inbound traffic & organic sign-ups



## Output-Driven Growth (Zero-CAC Loop)

### A. What It Does

- Embeds subtle NEXUS branding in every exported insight
- Distributes outputs via Slack, Google Docs, PDFs, and email

### B. How It Delivers Impact

- Recipients experience superior structured intelligence
- Creates a natural “Wow → Try → Adopt” loop
- Sustains perpetual, zero-cost user acquisition



## Enterprise Rollout with Data & Compliance Hooks

### A. What It Does

- Provides enterprise security: SOC2, ISO, RBAC, SSO
- Supports VPC / On-Prem deployments

### B. How It Delivers Impact

- Enables adoption across regulated or high-compliance teams
- Ensures sensitive research never leaves the organization
- Automates policy, compliance, and governance workflows at scale



CORE SUBSCRIPTION REVENUE	PLATFORM REVENUE STREAMS	ENTERPRISE & UPSELL REVENUE
PRO Plan – \$49/month or \$470/year	Workflow Marketplace (15% revenue share)	White-Label Licensing (\$150K-\$1M/yr)
Unlimited agent tasks; Advanced context graph; Custom workflows; Priority execution	Vertical workflow packs; Community workflows; Marketplace flywheel	Private deployments; Internal workflow libraries; Secure environments
Enterprise Plan – \$39–\$89/user/month	Intelligence & Data Products (anonymized trend insights)	Automation Packs (competitor monitors, regulatory alerts, exec briefings)
SSO, RBAC, audit trails, admin console	Sector heatmaps; Research dashboards; Premium briefings	Enterprise dataset indexing; Org-wide federated search
VPC/On-Prem deployment; SLA support	API & Developer Platform (orchestration & workflow APIs)	Deep enterprise penetration & high switching cost



PITFALLS	DESCRIPTION	MITIGATION STRATEGY
1. Multi-Agent Reasoning Drift ● High-Severity	Model degradation reduces synthesis accuracy and produces inconsistent insights.	Continuous fine-tuning; RLHF; Confidence scoring fallback LoE: ● Low
2. SEO/Ad-Biased Source Pollution ● High-Severity	Sponsored/SEO pages outrank authoritative sources causing polluted retrieval.	Relevance Integrity Layer; Multi-source validation; Domain trust scoring ,LoE: ● Medium
3. Context Graph Noise ● Medium-Severity	Graph accumulates outdated/irrelevant nodes reducing resurfacing precision.	Graph pruning; Embedding recalibration; User tagging LoE: ● Low
4. Workflow Automation Breakage ● High-Severity	External site/API changes break triggers causing workflow failures.	Redundant triggers; Auto-retry + anomaly detection; Health monitoring , LoE: ● Medium
5. Cross-Agent Output Inconsistency ● Medium-Severity	Parallel agents produce conflicting results leading to synthesis errors.	Contradiction scoring; Majority voting; Rule-based synthesis agent LoE: ● Low
6. Data Privacy Risks ● Medium -Severity	Sensitive enterprise research risks compliance exposure.	SOC2 encryption; Zero-retention mode; VPC/on-prem deployment LoE: ● Medium

## FUTURE ENHANCEMENTS

### Adaptive Agent Ensembles (Dynamic Orchestration Layer)

- Multi-agent teams automatically reorganize based on task complexity, domain, and prior performance.
- → Increases insight accuracy by 20–30% and reduces reasoning drift.

### Domain-Specific Intelligence Packs

- Pretrained domain packs for Finance, Technology with specialized retrieval, terminology understanding, and workflow templates.
- → Reduces onboarding time and improves relevance for vertical users.

## SUCCESS METRICS

### 1. North Star Metric – Verified Multi-Agent Insight Accuracy

$$\text{InsightAccuracy} = \frac{\text{InsightsValidatedBy} \geq 3\text{Sources}}{\text{TotalInsightsgenerated}} \times 100$$

- Purpose: Measures the reliability of synthesized insights.
- Goal: ≥ 92% Verified Insight Accuracy
- Tracks correctness, cross-source triangulation, and bias reduction.

### 2. Growth Metric – Active Workflows Per Org

$$\text{OrgWorkflowGrowth} = \frac{(\text{ActiveWorkflows} + \text{NewTeamsActivated} - \text{ChurnedTeams})}{\text{TotalTeams}} \times 100$$

- Purpose: Measures adoption + productivity impact.
- Goal: +60% quarterly team activation
- Indicates team-level network effects and PLG expansion.

### 3. Source Integrity Index

$$\text{SourceIntegrity} = \frac{(\text{VerifiedSources} - \text{SEO/SponsoredSources})}{\text{TotalSourcesUsed}} \times 100$$

- Purpose: Tracks how well NEXUS filters out SEO bias, sponsored content, misinfo.
- Goal: ≥ 90% authoritative sources
- Ensures clean retrieval pipelines.

### 4. Retention Efficiency (AARRR)

$$\text{RetentionEfficiency} = \frac{\text{DAU} + \text{WeeklyWorkflowRuns}}{\text{MAU}} \times 100$$

- Purpose: Measures habit formation and automated workflow stickiness.
- Goal: 40–55% DAU/MAU ratio
- Matches top-tier PLG benchmarks.

[Home](#)[Services](#)[Features](#)[About Us](#)[Contact](#)

# NEMUS

MULTI-AGENT BROWSER

TRANSFORMING KNOWLEDGE WORK THROUGH ORCHESTRATED INTELLIGENCE

# MAIN DASHBOARD

## NEXUS

Multi-Agent Research OS

- [Home](#)
  - [Agents](#)
  - [Workflows](#)
  - [Memory Graph](#)
  - [Alerts](#)
  - [Projects](#)
  - [Settings](#)
- System Status: Active

## Welcome back

Your intelligent research workspace is ready

Ask NEXUS to research anything...

Try: "Compare 3 fintech competitors" or "Deep dive on Series B funding"

## Quick Actions

- [!\[\]\(762e78a51e928061c95df2864acb5b1c\_img.jpg\) Competitive Landscape  
Analyze market competitors and positioning](#)
- [!\[\]\(dcd9ac94df8c90ed63ada777f031e1bd\_img.jpg\) Funding Deep Dive  
Research funding rounds and investors](#)
- [!\[\]\(f29e10600d51605a8e57ed20669457a7\_img.jpg\) Market Scan  
Monitor industry trends and signals](#)
- [!\[\]\(ed97ce7c85186184abf4d9253dc670f9\_img.jpg\) Regulatory Tracker  
Track compliance and policy changes](#)

## Active Agents

- [!\[\]\(013778809bd67eadafb80735b3bce3e4\_img.jpg\) Retrieval Agent  
active  
Fetching source #12...  
Evaluating credibility...  
Processing data stream...](#)
- [!\[\]\(10cd921ffce1ed4b4b54462932629d27\_img.jpg\) Verification Agent  
active  
Cross-validating claims...  
Checking source reliability...  
Analyzing patterns...](#)
- [!\[\]\(1c5d36d6ded01242d3cbf8e1c0865272\_img.jpg\) Synthesis Agent  
idle  
Waiting for data...  
Ready to synthesize...  
Standing by...](#)
- [!\[\]\(9aa11ef272e6b332eadcc794da5fa480\_img.jpg\) Memory Agent  
complete  
Graph updated  
Entities mapped  
Task complete](#)

## Recent Projects

- [!\[\]\(96e9283c6e911eb1923e9921ebea56d6\_img.jpg\) Fintech Market Analysis  
Complete • 2 hours ago](#)
- [!\[\]\(bc595644e1940818ac3bb98c94dd9e45\_img.jpg\) Series B Investor Mapping  
In Progress • 1 day ago](#)
- [!\[\]\(ed3310508e67c104b53327e7daf687b9\_img.jpg\) AI Regulatory Landscape  
Complete • 3 days ago](#)

### Universal Command Bar

Natural-language interface to trigger any research workflow instantly ("Compare top fintech competitors...").

### Quick Action Templates

Pre-built workflows for high-frequency tasks like Competitive Analysis, Funding Deep Dive, Market Scans, or Regulatory Tracking.

### Active Agents Panel

Real-time view of the multi-agent engine – retrieval, verification, synthesis, ranking, and memory agents all working in parallel.

# SEARCH RESULTS PAGE WITH AI-POWERED FEATURES.

**NEXUS**  
Multi-Agent Research OS

- Home
- Agents
- Workflows
- Memory Graph
- Alerts
- Projects
- Settings

System Status • Active

## Analysis Setup

Configure your multi-agent research workflow

**Research Query**

Compare top 5 fintech competitors by funding and market positioning

**Region**: Global | **Time Window**: Last 12 Months | **Industry**: Fintech

**Example Queries**

- Compare top 5 fintech competitors by funding and market positioning
- Create market map for SaaS security startups with growth trends
- Generate investor mapping and funding analysis for Series B fintech startups
- Analyze regulatory compliance requirements across EU and US markets

**Run Multi-Agent Analysis**

**Research Query Input**  
Accepts natural-language queries which NEXUS decomposes into multi-step tasks.

**Context Controls**

- Region: Global, APAC, US, EU
- Time Window: Last 1–12 months
- Industry: Fintech, SaaS, Healthcare, etc.
- Enables hyper-targeted research.

**Run Multi-Agent Analysis**  
Launches Retrieval, Verification, Memory, Synthesis, and Ranking agents in parallel.

# MULTI-AGENT ANALYSIS WIREFRAME

NEXUS

Multi-Agent Research OS

## Multi-Agent Analysis

Competitive analysis

Home

Agents

Workflows

Memory Graph

Alerts

Projects

Settings

System Status

● Active

### Overall Progress

21%

#### Retrieval Agent

active

Retrieved 156 data points  
Filtering by relevance...  
Download complete

#### Verification Agent

active

Cross-checking revenue figures...  
Validating source credibility...  
3 conflicting data points detected...

#### Synthesis Agent

idle

Waiting for verification...  
Ready to synthesize...

#### Memory Agent

idle

Preparing graph connections...  
Standing by...

#### Ranking Agent

idle

Awaiting synthesis...  
Ready to rank...

#### Report Agent

idle

Awaiting final data...  
Ready to compile...

### Overall Progress Indicator

Aggregates completion of all agents – retrieval → verification → synthesis → ranking → reporting.

- Retrieval Agent: Gathers, filters, and pre-processes data.
- Verification Agent: Cross-checks claims, resolves inconsistencies, detects SEO-bias.
- Synthesis Agent: Generates structured insights.
- Memory Agent: Links new findings to past analyses.
- Ranking Agent: Orders insights based on relevance, credibility, and context.
- Report Agent: Prepares export-ready structured outputs.

# MEMORY GRAPH (CONTEXTUAL INTELLIGENCE LAYER)

## NEXUS

Multi-Agent Research OS

- Home
- Agents
- Workflows
- Memory Graph**
- Alerts
- Projects
- Settings

System Status: Active

Search nodes...

🔍✖⟳

The graph displays entities as nodes and their relationships as edges. Nodes include Stripe (blue), Revolut (blue), Visa (blue), Plaid (blue), Fintech Fund (green), and Series B Trends (green). Relationships shown are: Stripe and Revolut are connected by 'Recent Round'; Stripe and Visa are connected by 'Partnership'; Stripe is connected to 'Shows' (Metadata, Linked insights, Related documents); Stripe and Fintech Fund are connected by 'Market Leader'; Fintech Fund and Plaid are connected by 'Data Provider'.

**Stripe**

Type: company  
Connections: 2  
Last updated: 2 hours ago

**Linked Insights**

- Fintech Market Analysis (Dec 2024)
- Payment Processing Trends

**Related Documents**

- Series B Funding Report.pdf
- Competitive Analysis Q4 2024

**Actions:** Pin Node, Add Note

## KEY FEATURES

### Entity Nodes & Relationships

Each node (company, topic, trend) is linked with meaningful relationships:

### Dynamic Visualization

Nodes scale and cluster based on relevance, confidence, and recency.

### Searchable Knowledge Base

Users can instantly jump to any past insight, entity, or dataset.

### Persistent Memory

Insights from every analysis automatically update the graph, creating enterprise-grade domain intelligence over time.

# ALERTS & MONITORING WIREFRAMES

## NEXUS

Multi-Agent Research OS

- Home
- Agents
- Workflows
- Memory Graph
- Alerts
- Projects
- Settings

System Status: Active

### Alerts & Monitoring

Automated intelligence tracking for your research workflows

2 Unread Alerts New

5 Monitoring Workflows Active

07:00 Next Scan Tomorrow Scheduled

**Stripe raised \$600M Series H** new

△ Funding Event Confidence: 96% 2 hours ago

**Visa announced partnership with PayPal** new

△ Partnership Confidence: 89% 5 hours ago

**EU PSD3 draft regulations released** reviewed

△ Regulatory Confidence: 94% 1 day ago

**Plaid user base grew 34% QoQ** reviewed

△ Growth Metric Confidence: 82% 2 days ago

**Automated Monitoring Workflows:**

- NEXUS runs scheduled multi-agent scans (news, filings, regulatory portals, datasets) at predefined intervals.
- Autonomous retrieval + verification ensures only high-confidence updates are surfaced.

**Entity-Aware Tracking**

- Alerts automatically link to related entities in the Memory Graph.
- Example: "Stripe raised \$600M Series H" → ties to Stripe node + Funding Trends cluster.
- Enables instant contextual expansion from alert → deeper analysis.